


# Multi-layer Perceptron

Artificial Intelligence


School of Computer Science  
The University of Adelaide

# Recall Perceptron




**Perceptron:** A neural network with only a single layer. Each input has a connection to every neuron in the network.

# Recall Perceptron




XOR




Single-layer perceptron is a linear classifier.

**Perceptron can not solve non-linear problems.**

# Multi-layer Perceptron - A Non-linear Classifier



**Perceptron**




**MLP**

MLPs are more expressive than Perceptrons since they can learn highly non-linear class boundaries.


# Multi-layer Perceptron - A Non-linear Classifier

The most common case involves a single hidden layer:




# Multi-layer Perceptron - A Non-linear Classifier

Each *hidden unit* can be considered as single output perceptron network:



# Multi-layer Perceptron - A Non-linear Classifier


Which is capable of separating the training examples linearly



# Multi-layer Perceptron




## - A Non-linear Classifier

The output unit of the multi-layer network can then be considered a soft-thresholded linear combination of the hidden units (which are equivalent to single output unit perceptrons):






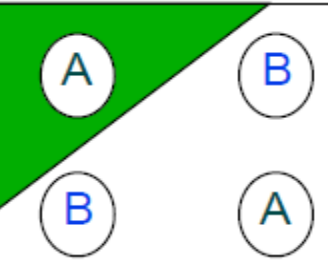


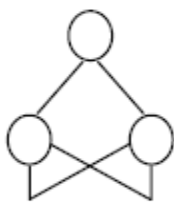
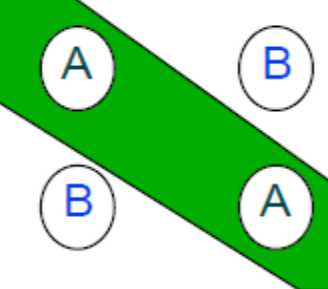
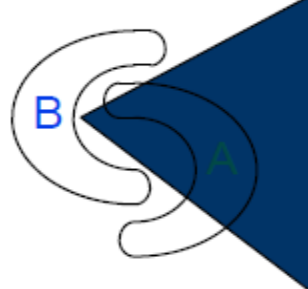

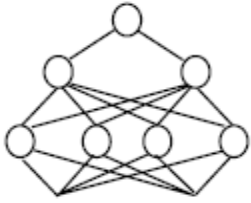
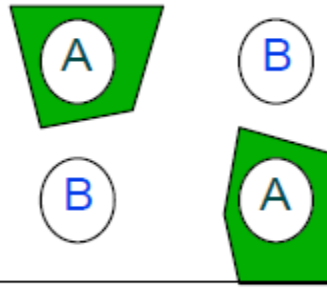

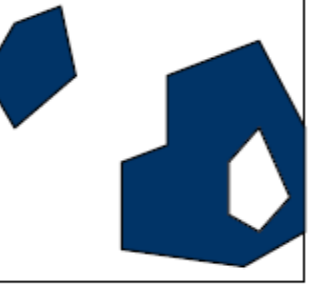
# Multi-layer Perceptron - A Non-linear Classifier



- (a) The result of combining two opposite-facing soft threshold functions to produce a ridge.
- (b) The result of combining two ridges to produce a bump.

# Multi-layer Perceptron

## A Nonlinear Classifier

Structure	Types of Decision Regions	Exclusive-OR Problem	Classes with Meshed regions	Most General Region Shapes
<b>Single-Layer</b> 	<b>Half Plane Bounded By Hyperplane</b>			
<b>Two-Layer</b> 	<b>Convex Open Or Closed Regions</b>			
<b>Three-Layer</b> 	<b>Arbitrary (Complexity Limited by No. of Nodes)</b>			

# Multi-layer Perceptron


## The Good and the Bad

- With a single, sufficiently large hidden layer, it is possible to represent *any continuous* function of the inputs with *arbitrary* accuracy.
- Unfortunately, for any *particular* network, it is harder to characterize exactly which functions can be represented and which ones cannot.
- As a consequence, given a particular learning problem, it is unknown how to choose the *right number of hidden units* in advance.
- One usually resorts to cross validation, but this can be computationally expensive for large networks.

# Training MLP

- Backpropagation


# Gradient Descent



**Figure 5.3** The first step in iteratively finding the minimum of this loss function, by moving  $w$  in the reverse direction from the slope of the function. Since the slope is negative, we need to move  $w$  in a positive direction, to the right. Here superscripts are used for learning steps, so  $w^1$  means the initial value of  $w$  (which is 0),  $w^2$  at the second step, and so on.

Weight change: 
$$w^{t+1} = w^t - \eta \frac{d}{dw} f(x; w)$$


# Gradient Descent






# Learning Rate


Big learning rate



Small learning rate



# Understanding Perceptron Training



- The function that a perceptron network corresponds to can be represented as  $h_{\mathbf{W}}(\mathbf{x})$ , where

$$h_{\mathbf{W}}(\mathbf{x}) = g(inputs) = g\left(\sum_{j=0}^D W_j x_j\right)$$



# Understanding Perceptron Training

- Perceptron learning (generally, neural network learning) occurs by *adjusting the weights to minimize some measure of error*.
- Let  $(\mathbf{x}, y)$  be a *single* training sample with its *true* output  $y$ . The squared error is given by

$$\begin{aligned} E &= \frac{1}{2} Err^2 \\ &= \frac{1}{2} (y - h_{\mathbf{W}}(\mathbf{x}))^2 \\ &= \frac{1}{2} (y - g(\sum_{j=0}^D W_j x_j))^2 \end{aligned}$$

Note scaling the error with  $\frac{1}{2}$  does not change its minimizer.

# Understanding Perceptron Training

- Calculating the partial derivative of the error against a particular weight, we have

$$\begin{aligned}\frac{\partial E}{\partial W_j} &= Err \times \frac{\partial Err}{\partial W_j} \\ &= Err \times \frac{\partial}{\partial W_j} \left( y - g\left(\sum_{j=0}^D W_j x_j\right) \right) \\ &= -Err \times g'\left(\sum_{j=0}^D W_j x_j\right) \times x_j\end{aligned}$$

where  $g'$  is the derivative of the activation function  $g$ .

# Understanding Perceptron Training

- Under the **gradient descent** algorithm, if we want to *reduce*  $E$ , we update the weight as follows:

$$W_j \longleftarrow W_j + \alpha \times Errr \times g'(\sum_{j=0}^D W_j x_j) \times x_j$$

where  $\alpha$  is the **learning rate**.


# Backpropagation

## Multi-output Multi-layer Perceptron

- We need to consider multiple output units for multi-layer networks. Let  $(\mathbf{x}, \mathbf{y})$  be a single sample with its desired output labels  $\mathbf{y} = \{y_1, \dots, y_i, \dots, y_M\}$ .
- The error at the output units is just  $\mathbf{y} - h_{\mathbf{W}}(\mathbf{x})$ , and we can use this to adjust the weights between the hidden layer and the output layer.
- The above steps produces a term equivalent to the error at the hidden layer, i.e. the error at the output layer is **back-propagated** to the hidden later.
- This is subsequently used to update the weights between the input units and the hidden layer.

# Backpropagation

## Muti-output Multi-layer Perceptron





# Backpropagation

Step 1: Update the weights between the hidden and output layers.

- Let  $Err_i$  be the  $i$ -th component of the error vector  $\mathbf{y} - h_{\mathbf{W}}(\mathbf{x})$ .
- Define  $\Delta_i = Err_i \times g'(in_i)$ .
- The weight-update rule becomes

$$W_{j,i} \longleftarrow W_{j,i} + \alpha \times a_j \times \Delta_i$$

This is similar to weight-updates for Perceptrons!

# Backpropagation

Step 2: Back-propagate the error to the hidden layer.

- The idea is that the hidden node  $j$  is “responsible” for some fraction of the error  $\Delta_i$  in each of the output nodes to which it connects.
- Thus the  $\Delta_i$  values are divided according to the strength (weight) of the connection between the hidden node and the output node:

$$\Delta_j = g'(in_j) \sum_i W_{j,i} \Delta_i$$

# Backpropagation

Step 3: Update the weights between the input units and the hidden layer.

- Again, this is similar to weight-updates in Perceptrons:

$$W_{k,j} \longrightarrow W_{k,j} + \alpha \times a_k \times \Delta_j$$



# Backpropagation

For the general case of *multiple hidden layers*:

- ① Compute the  $\Delta$  values for the output units, using the observed error.
- ② Starting with the output layer, repeat the following for each layer in the network, until the earliest hidden layer is reached:
  - Propagate the  $\Delta$  values back to the previous layer.
  - Update the weights between the two layers.
- ③ Repeat Steps 1 to 2 for all training samples.

# Backpropagation

**function** BACK-PROP-LEARNING(*examples*, *network*) **returns** a neural network

**inputs:** *examples*, a set of examples, each with input vector  $\mathbf{x}$  and output vector  $\mathbf{y}$   
*network*, a multilayer network with  $M$  layers, weights  $W_{j,i}$ , activation function  $g$

**repeat**

**for each**  $e$  **in** *examples* **do**

**for each** node  $j$  in the input layer **do**  $a_j \leftarrow x_j[e]$  Initialize

**for**  $\ell = 2$  **to**  $M$  **do** forward

$in_i \leftarrow \sum_j W_{j,i} a_j$

$a_i \leftarrow g(in_i)$

**for each** node  $i$  in the output layer **do** backward

$\Delta_i \leftarrow g'(in_i) \times (y_i[e] - a_i)$

**for**  $\ell = M - 1$  **to**  $1$  **do**

**for each** node  $j$  in layer  $\ell$  **do**

$\Delta_j \leftarrow g'(in_j) \sum_i W_{j,i} \Delta_i$


**for each** node  $i$  in layer  $\ell + 1$  **do**

$W_{j,i} \leftarrow W_{j,i} + \alpha \times a_j \times \Delta_i$

**until** some stopping criterion is satisfied

**return** NEURAL-NET-HYPOTHESIS(*network*)

# Backpropagation Derivation



$$E = \frac{1}{2} \sum_i (y_i - a_i)^2$$

$$\Delta_i = Err_i \times g'(in_i)$$



$$\Delta_j = g'(in_j) \sum_i W_{j,i} \Delta_i$$

# Backpropagation

## Multi-layer Perceptron

- Chain rule:

Given  $f(x) = u(v(x))$ , the derivative of  $f(x)$  respect to  $x$  is:


$$\frac{df}{dx} = \frac{du}{dv} \cdot \frac{dv}{dx}$$

Given  $f(x) = u(v(w(x)))$ , the derivative of  $f(x)$  respect to  $x$  is:

$$\frac{df}{dx} = \frac{du}{dv} \cdot \frac{dv}{dw} \cdot \frac{dw}{dx}$$

# Backpropagation with SGD

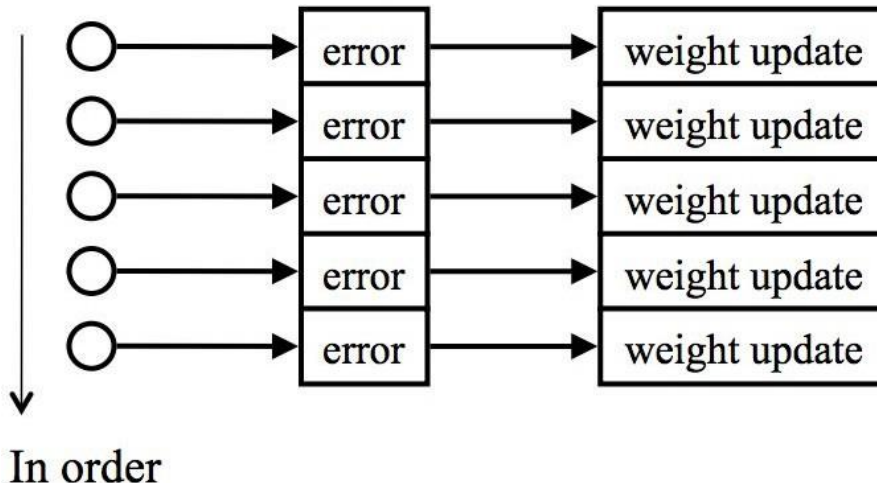
Batch: training step



## Gradient Decent

**Batch:** training over all given examples once.



Online: training step




## Stochastic Gradient Decent

- Randomly choose  $m$  samples
- Compute the gradients
- Do backpropagation
- Repeat

# Deep Multi-layer Perceptron

- We can learn anything !!!
- More than one hidden layer  **deep.**
- Higher level representations  Better at high-level tasks.
- Visual Classification / Speech Recognition / Scene understanding / Visual question answering ....



# Not so fast... ☹️

- Backpropagation [late 80s, early 90s]
  - Goal was to train nets with large number of layers, so that features could be learned directly from input data, but it didn't quite work
  - Notable exception: convolutional neural net by Y. LeCun (large # layers, but small # parameters)
- Issues with MLP training via backpropagation
  - Very slow convergence, particularly in large nets and large databases
  - Slow computers of the 80s and 90s
  - Local minima (how to initialize SGD)
  - Net structure (cross validation)
  - Overfitting