Approximate Inference

Artificial Intelligence

School of Computer Science
The University of Adelaide

Inference on Bayesian Networks

Exact inference: computational expensive for a large BN.

- Number of multiplications approach to $O(n2^n)$

$$\begin{aligned} P(b|j,\neg m) &= \alpha \sum_e \sum_a P(b)P(e)P(a|b,e)P(j|a)P(\neg m|a) \\ &= \alpha\,[P(b)P(e)P(a|b,e)P(j|a)P(\neg m|a) \\ &\quad +P(b)P(e)P(\neg a|b,e)P(j|\neg a)P(\neg m|\neg a) \\ &\quad +P(b)P(\neg e)P(a|b,\neg e)P(j|a)P(\neg m|a) \\ &\quad +P(b)P(\neg e)P(\neg a|b,\neg e)P(j|\neg a)P(\neg m|\neg a)] \end{aligned}$$

Approximate inference:

- Approximately calculate the posterior probability.
- Use random sampling for inference.
- More samples leads to more accurate solutions.

Inference on Bayesian Networks

Exact inference: computational expensive for a large BN.

- Number of multiplications approach to $O(n2^n)$

Approximate inference with sampling:

Also called Monte Carlo Methods.

- Direct Sampling (Prior Sampling).
- Rejection Sampling.
- Likelihood Weighting.
- Gibbs Sampling

Direct sampling

Markov chain sampling

Sampling

- What is sampling

  Sampling is a statistical procedure to select the individual observations from the population.

- Why sampling

  Statisticians attempt for the samples to represent the whole population in question.

- Example:

  What is the probability of getting 3 when rolling a dice?

$$P(X = x_i) = \frac{number\ of\ times\ \{X = x_i\}}{total\ number\ of\ trials}$$

Sampling from a Distribution

Sample from a given distribution of a Variable.

- Given the uniform distribution of random variable *W.*

  values: $\{w_1, w_2, \ldots w_n\}$,

  corresponding probabilities: $p_1, p_2, \ldots, p_n$, $\sum_i p_i = 1$.

  - Get a sample *u* from uniform distribution in [0,1).
    In python : random()
  - Map *u* to a specific instantiation of *W*.

$$
\begin{array}{rl}
0 \leq u < p_1 & w_1 \\
p_1 \leq u < p_1 + p_2 & w_2 \\
p_1 + p_2 \leq u < p_1 + p_2 + p_3 & w_3 \\
p_1 + p_2 + \ldots + p_{n-1} \leq u < p_1 + p_2 + \ldots + p_{n-1} + p_n = 1 & w_n
\end{array}
$$

Sampling from a Distribution

Example of sampling from a given distribution of a Variable.

- Sample space of W: {*Sunny, Rain, Cold, Snow*}
- W follows uniform distribution

Weather (*W*)	**P(*W* = w)**
Sunny	**0.25**
Rain	**0.25**
Cold	**0.25**
Snow	**0.25**

$$
\begin{aligned}
0 \leq u < 0.25 &\Rightarrow W = sunny \\
0.25 \leq u < 0.5 &\Rightarrow W = rain \\
0.5 \leq u < 0.75 &\Rightarrow W = cold \\
0.75 \leq u < 1 &\Rightarrow W = snow
\end{aligned}
$$

$$P(X = x_i) = \frac{number\ of\ times\ \{X = x_i\}}{total\ number\ of\ trials}$$

Sampling from a Distribution

Example of sampling from a given distribution of a Variable.

- Sample space of W: {*Sunny, Rain, Cold, Snow*}
- W follows the distribution provided in the table.

Weather (*W*)	**P(*W* = w)**
Sunny	**0.3**
Rain	**0.3**
Cold	**0.3**
Snow	**0.1**

$$
\begin{aligned}
0.0 \leq u < 0.3 &\Rightarrow W = sunny \\
0.3 \leq u < 0.6 &\Rightarrow W = rain \\
0.6 \leq u < 0.9 &\Rightarrow W = cold \\
0.9 \leq u < 1.0 &\Rightarrow W = snow
\end{aligned}
$$

Direct Sampling

How to sample from a given distribution of Variables in BN?

- Sample each variable in turn, in a topological order of a BN.
- Previously sampled variables value are used as condition for sampling the next node of BN.

Direct Sampling

How to sample from a given distribution of Variables in BN?

- Sample each variable in turn, in a topological order of a BN.
- Previously sampled variables value are used as condition for sampling the next node of BN.

Example: Wet Grass Network.

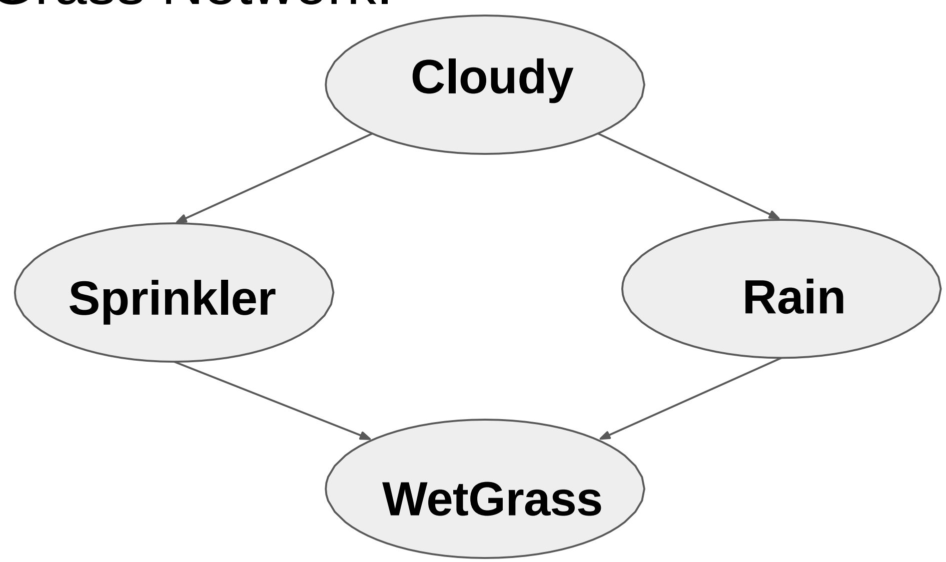

Direct Sampling

$P(C)$

P(c)
0.5

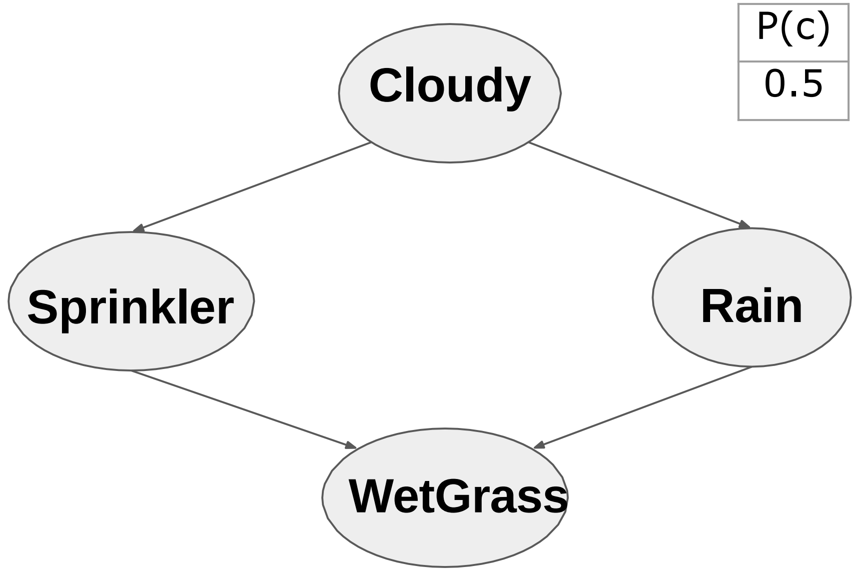

$P(S|C)$

C	P(s|C)
c	0.1
-c	0.5

$P(R|C)$

C	P(r|C)
c	0.8
-c	0.2

$P(W|S,R)$

S	R	P(w|S,R)
s	r	0.99
s	-r	0.90
-s	r	0.90
-s	-r	0.01

Direct Sampling

$P(C)$

c	0.5
-c	0.5

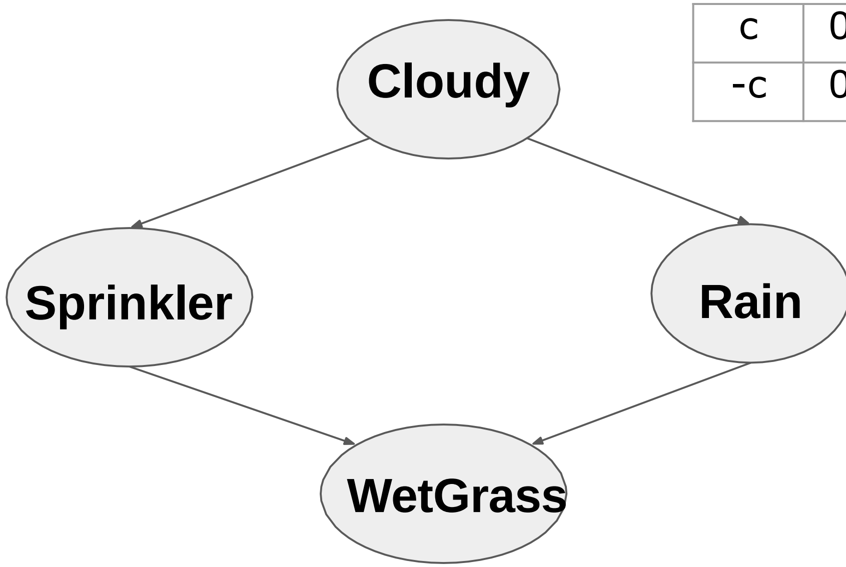

$P(S|C)$

c	s	0.1
	-s	0.9
-c	s	0.5
	-s	0.5

$P(R|C)$

c	r	0.8
	-r	0.2
-c	r	0.2
	-r	0.8

$P(W|S,R)$

s	r	w	0.99
		-w	0.01
	-r	w	0.90
		-w	0.10
-s	r	w	0.90
		-w	0.10
	-r	w	0.01
		-w	0.99

Direct Sampling

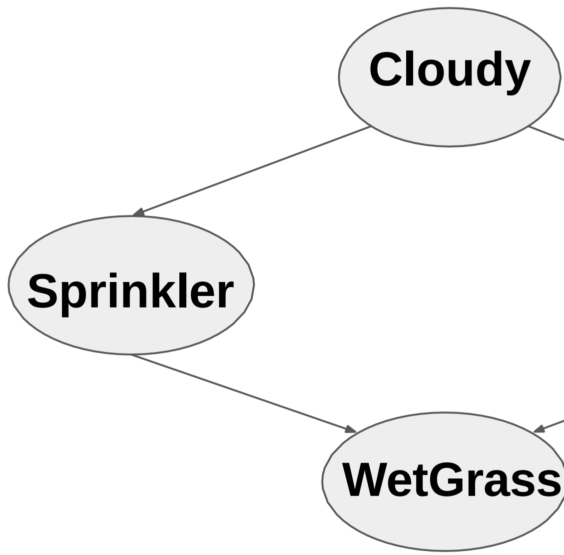

$P(C)$

c	0.5
-c	0.5

$P(S|C)$

c	s	0.1
	-s	0.9
-c	s	0.5
	-s	0.5

$P(R|C)$

c	r	0.8
	-r	0.2
-c	r	0.2
	-r	0.8

$P(W|S,R)$

s	r	w	0.99
		-w	0.01
	-r	w	0.90
		-w	0.10
-s	r	w	0.90
		-w	0.10
	-r	w	0.01
		-w	0.99

- **Fix sampling order**
  - **{C, S, R,W}**
- **Sample examples given CPTs.**

Direct Sampling

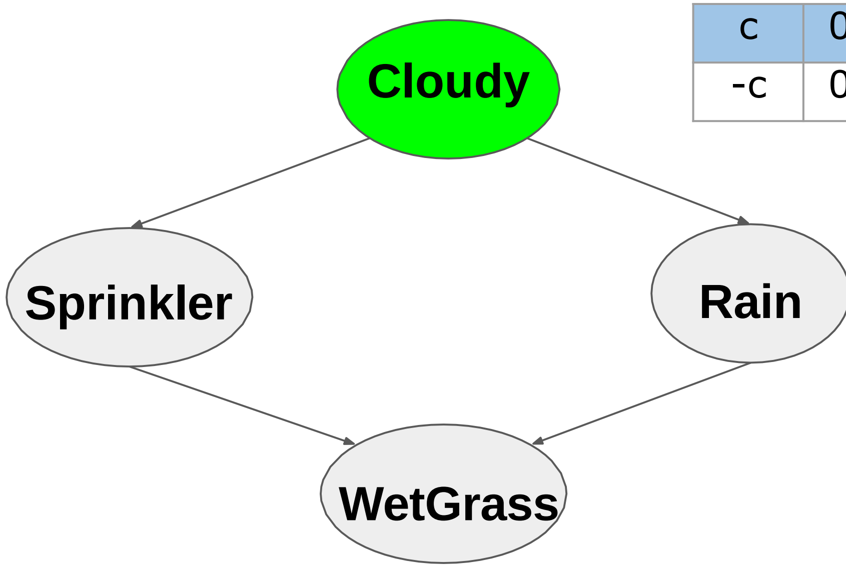

$P(C)$

c	0.5
-c	0.5

$P(S|C)$

c	s	0.1
	-s	0.9
-c	s	0.5
	-s	0.5

$P(R|C)$

c	r	0.8
	-r	0.2
-c	r	0.2
	-r	0.8

$P(W|S,R)$

s	r	w	0.99
		-w	0.01
	-r	w	0.90
		-w	0.10
-s	r	w	0.90
		-w	0.10
	-r	w	0.01
		-w	0.99

$u = 0.22$

- **Fix sampling order**
  - **{C, S, R,W}**
- **Sample examples given CPTs.**
  - **{+c,**

Direct Sampling

$P(C)$

c	0.5
-c	0.5

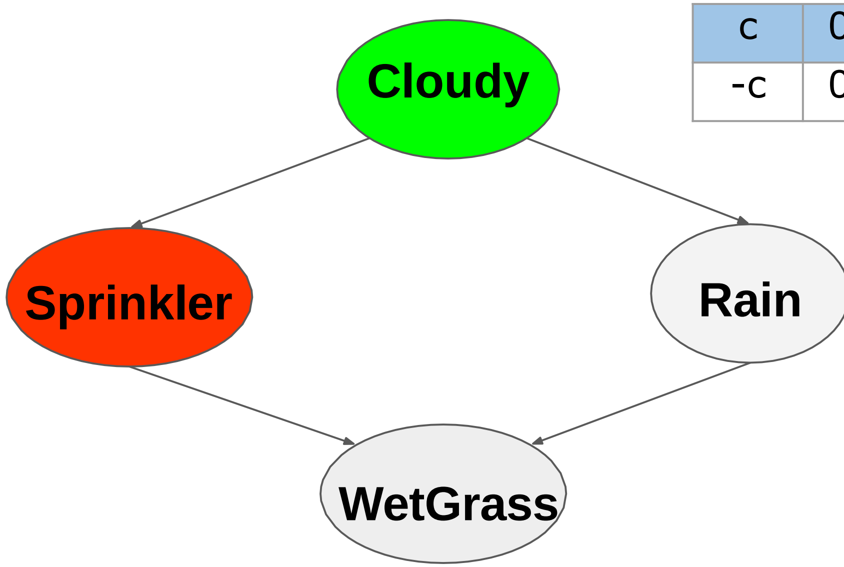

$P(S|C)$

c	s	0.1
	-s	0.9
-c	s	0.5
	-s	0.5

$P(R|C)$

c	r	0.8
	-r	0.2
-c	r	0.2
	-r	0.8

$P(W|S,R)$

s	r	w	0.99
		-w	0.01
	-r	w	0.90
		-w	0.10
-s	r	w	0.90
		-w	0.10
	-r	w	0.01
		-w	0.99

$u = 0.81$

- **Fix sampling order**
  - **{C, S, R,W}**
- **Sample examples given CPTs.**
  - **{+c, -s,**

Direct Sampling

$P(C)$

c	0.5
-c	0.5

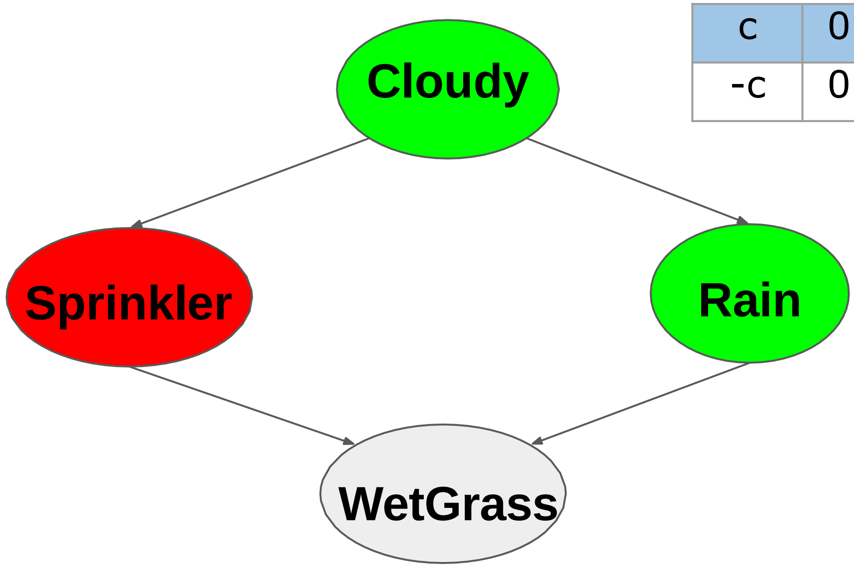

$P(S|C)$

c	s	0.1
	-s	0.9
-c	s	0.5
	-s	0.5

$P(R|C)$

c	r	0.8
	-r	0.2
-c	r	0.2
	-r	0.8

$P(W|S,R)$

s	r	w	0.99
		-w	0.01
	-r	w	0.90
		-w	0.10
-s	r	w	0.90
		-w	0.10
	-r	w	0.01
		-w	0.99

$u = 0.65$

- **Fix sampling order**
  - **{C, S, R,W}**
- **Sample examples given CPTs.**
  - **{+c, -s, +r,**

Direct Sampling

$P(C)$

c	0.5
-c	0.5

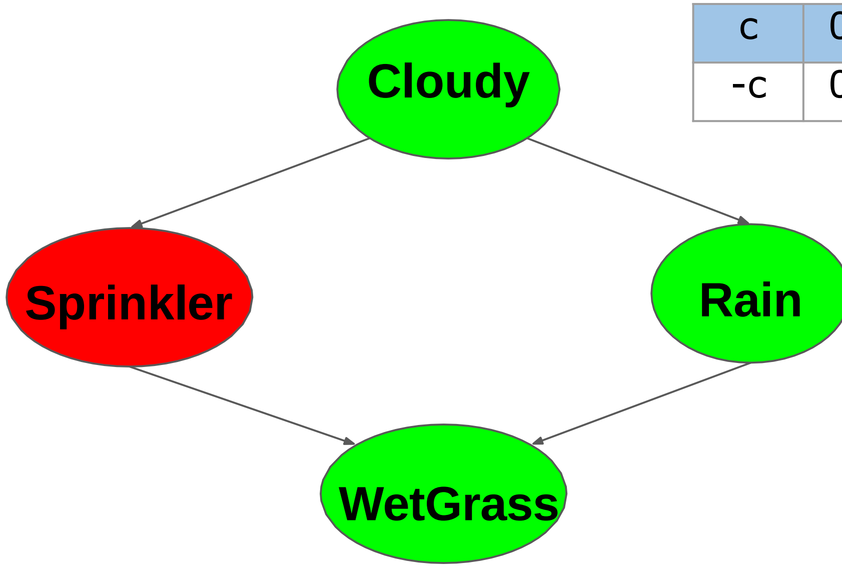

$P(S|C)$

c	s	0.1
	-s	0.9
-c	s	0.5
	-s	0.5

$P(R|C)$

c	r	0.8
	-r	0.2
-c	r	0.2
	-r	0.8

$P(W|S,R)$

s	r	w	0.99
		-w	0.01
	-r	w	0.90
		-w	0.10
-s	r	w	0.90
		-w	0.10
	-r	w	0.01
		-w	0.99

$u = 0.78$

- **Fix sampling order**
  - **{C, S, R,W}**
- **Sample examples given CPTs.**
  - **{+c, -s, +r, +w}**

Direct Sampling

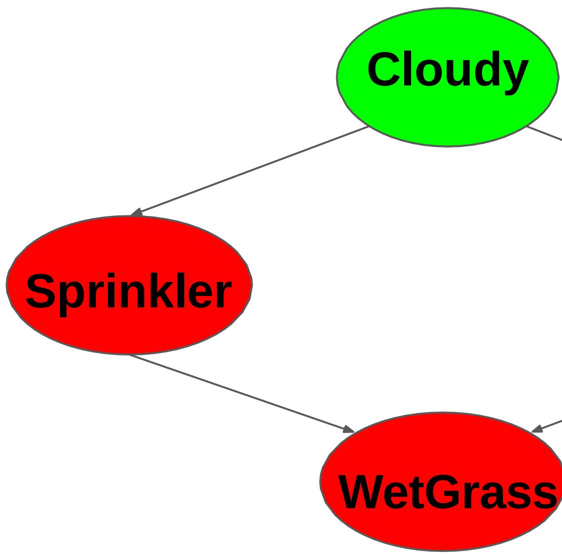

$P(C)$

c	0.5
-c	0.5

$P(S|C)$

c	s	0.1
	-s	0.9
-c	s	0.5
	-s	0.5

$P(R|C)$

c	r	0.8
	-r	0.2
-c	r	0.2
	-r	0.8

$P(W|S,R)$

s	r	w	0.99
		-w	0.01
	-r	w	0.90
		-w	0.10
-s	r	w	0.90
		-w	0.10
	-r	w	0.01
		-w	0.99

- **Fix sampling order**
  - **{C, S, R,W}**
- **Sample examples given CPTs.**
  - **{+c, -s, +r, +w}**
  - **{-c, +s, -r, +w}**
  - **….**
  - **....**
  - **{+c, -s, -r, -w}**

Direct Sampling

- Given *N* samples, and the number of samples for a specific event is $N_{PS}(x_1, \ldots, x_n)$, then approximate inference with sampling gives the probability of this event:

$$S_{PS}(\text{-c,+s,+r,-w}) = \lim_{N\to\infty} \frac{N_{PS}(-c,s,r,-w)}{N}$$

$$S_{PS}(\text{-c,+s,+r,-w}) \approx \frac{N_{PS}(-c,s,r,-w)}{N}$$

- Why direct sampling works?

The sampling process generates samples with following probability as each sampling step depends only on the parent values:

$$S_{PS}(x_1 \ldots x_n) = \prod_{i=1}^{n} P(x_i \,|\, parents(X_i)) = P(x_1 \ldots x_n)$$

Recall global semantics of the BN.

Direct Sampling

We frequently want to estimate the probability of partially specified events $P(x_1, \ldots, x_m)$ with $m < n$.

This can be approximated by

$$P(x_1, \ldots, x_m) \approx \frac{N_{PS}(x_1, \ldots, x_m)}{N}$$

where $N_{PS}(x_1, \ldots, x_m)$ is now the number of samples among N

Direct Sampling

- Given following set of samples:
  - {-c, +s, -r, +w}
  - {-c, -s, +r, +w}
  - {-c, +s, +r, -w}
  - {-c, +s, -r, +w}
  - {+c, +s, +r, +w}
  - {-c, -s, +r, +w}
  - {+c, +s, -r, -w}
  - {-c, +s, +r, -w}

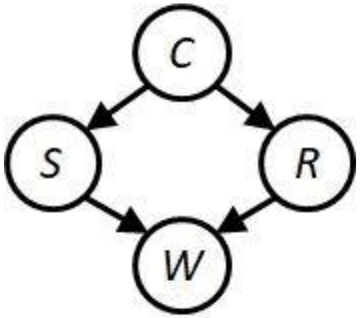

- P(-c, +s, +r, -w) = ?
- P(+r) = ?

Direct Sampling

- Given following set of samples:
  - {-c, +s, -r, +w}
  - {-c, -s, +r, +w}
  - {-c, +s, +r, -w}
  - {-c, +s, -r, +w}
  - {+c, +s, +r, +w}
  - {-c, -s, +r, +w}
  - {+c, +s, -r, -w}
  - {-c, +s, +r, -w}

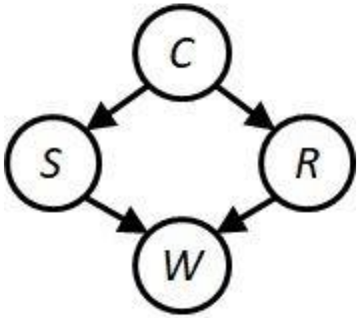

- P(-c, +s, +r, -w) = 2/8 = 0.25
- P(+r) = 5/8 = 0.625
- P(R | +w) = ?, P(-s | +w) = ?

Rejection Sampling

$$P(X|e)$$

- Generate samples as follows.
  - {-c, +s, -r, +w}
  - {-c, -s, +r, +w}
  - {-c, +s, +r, -w}
  - {-c, +s, -r, +w}
  - {+c, +s, +r, +w}
  - {-c, -s, +r, +w}
  - {+c, +s, -r, -w}
  - {-c, +s, +r, -w}

P(R | +w) = ?

P(-s | +w) = ?

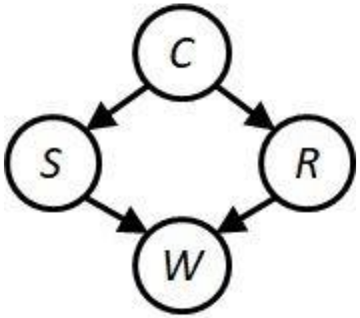

- Rejects the samples which does not match the evidence.

Rejection Sampling

$$P(X|e)$$

- Generate samples as follows.
  - {-c, +s, -r, +w}
  - {-c, -s, +r, +w}
  -
  - {-c, +s, -r, +w}
  - {+c, +s, +r, +w}
  - {-c, -s, +r, +w}
  -
  -

P(R | +w) = ?

P(-s | +w) = ?

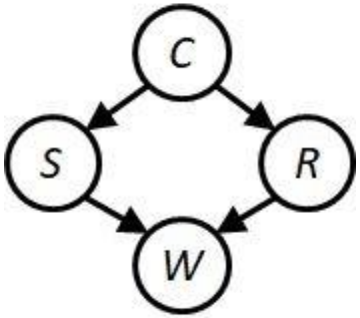

- Rejects the samples which does not match the evidence.
- $\hat{P}(X|e)$ is estimated by counting how many times X = x occurs for samples which are consistent with observations.

Rejection Sampling

$$P(X|e)$$

- Generate samples as follows.
  - {-c, +s, -r, +w}
  - {-c, -s, +r, +w}
  -
  - {-c, +s, -r, +w}
  - {+c, +s, +r, +w}
  - {-c, -s, +r, +w}
  -
  -

$$P(X|e) \approx \frac{N_{PS}(X,e)}{N_{PS}(e)}$$

P(R | +w) = <3/5,2/5>

P(-s | +w) = 2/5

- Rejects the samples which does not match the evidence.
- $\hat{P}(X|e)$ is estimated by counting how many times X = x occurs for samples which are consistent with observations.

Likelihood Weighting

- **Rejection sampling:**

  Inefficient with P(e) being small: We sample too many examples that are inconsistent with evidence.

- **Likelihood weighting**
  - Samples examples which are consistent with evidence.
  - Each sample have a support value w (i.e., weight).
    - Initialize w of the generated sample as w = 1
    - Repeat
      - If variable is non-evidence : sample as usual.
      - If variable is evidence variable E: set E = e, and set $w = w * P(E = e \mid parents(E))$

Example : Likelihood Weighting

Evidence: (+s,+w)

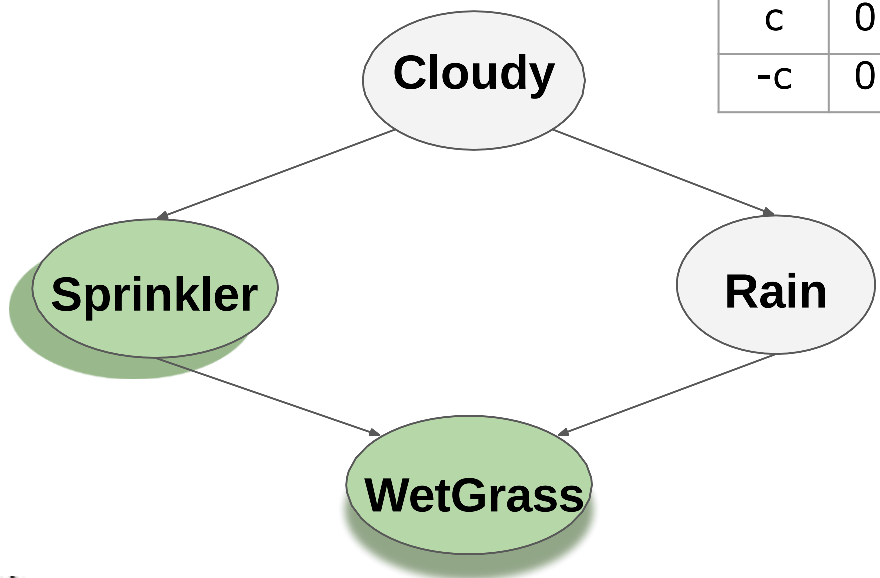

$P(C)$

c	0.5
-c	0.5

$P(S|C)$

c	s	0.1
	-s	0.9
-c	s	0.5
	-s	0.5

$P(R|C)$

c	r	0.8
	-r	0.2
-c	r	0.2
	-r	0.8

$P(W|S,R)$

s	r	w	0.99
		-w	0.01
	-r	w	0.90
		-w	0.10
-s	r	w	0.90
		-w	0.10
	-r	w	0.01
		-w	0.99

- **Fix sampling order {C, S, R, W}**
- **Likelihood weighted sampling with (+s, +w)**
  - **{_, _, _, _} , W = 1**

Example : Likelihood Weighting

Evidence: (+s,+w)

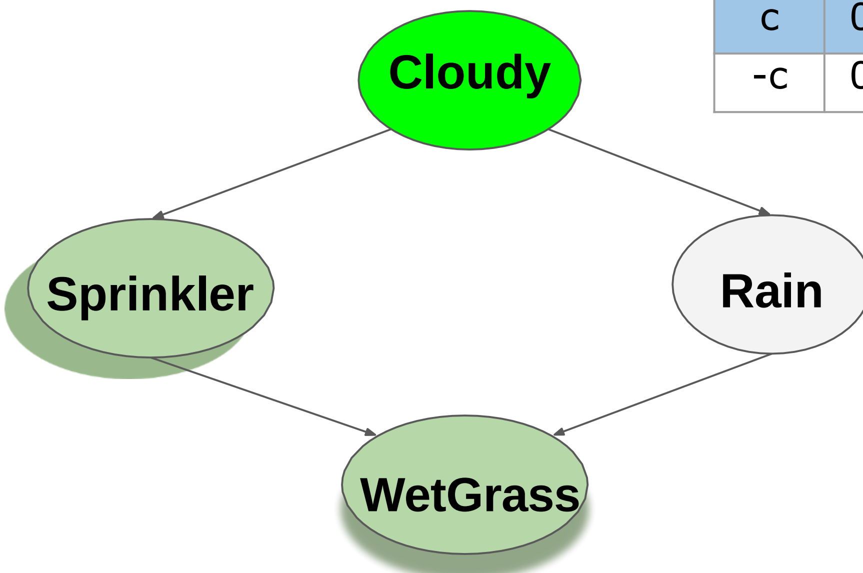

$P(C)$

c	0.5
-c	0.5

$P(S|C)$

c	s	0.1
	-s	0.9
-c	s	0.5
	-s	0.5

$P(R|C)$

c	r	0.8
	-r	0.2
-c	r	0.2
	-r	0.8

$P(W|S,R)$

s	r	w	0.99
		-w	0.01
	-r	w	0.90
		-w	0.10
-s	r	w	0.90
		-w	0.10
	-r	w	0.01
		-w	0.99

$u = 0.22$

- **Fix sampling order {C, S, R,W}**
- **Likelihood weighted sampling with (+s, +w)**
  - **{+c, _, _, _} , W = 1**

Example : Likelihood Weighting

Evidence: (+s,+w)

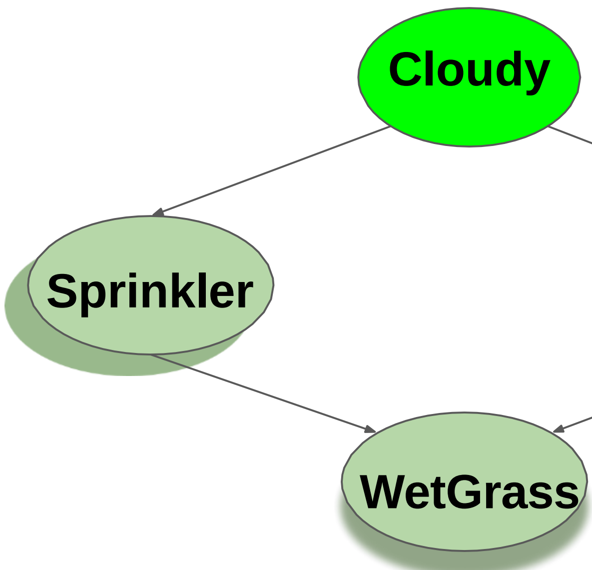

$P(C)$

c	0.5
-c	0.5

$P(S|C)$

c	s	0.1
	-s	0.9
-c	s	0.5
	-s	0.5

$P(R|C)$

c	r	0.8
	-r	0.2
-c	r	0.2
	-r	0.8

$P(W|S,R)$

s	r	w	0.99
		-w	0.01
	-r	w	0.90
		-w	0.10
-s	r	w	0.90
		-w	0.10
	-r	w	0.01
		-w	0.99

- **Fix sampling order {C, S, R,W}**
- **Likelihood weighted sampling with (+s, +w)**
  - **{+c, +s, _, _} , W = 1*0.1**

$P(+s|+c)$

$u = 0.81$
Not
Used!

Example : Likelihood Weighting

Evidence: (+s,+w)

$P(C)$

c	0.5
-c	0.5

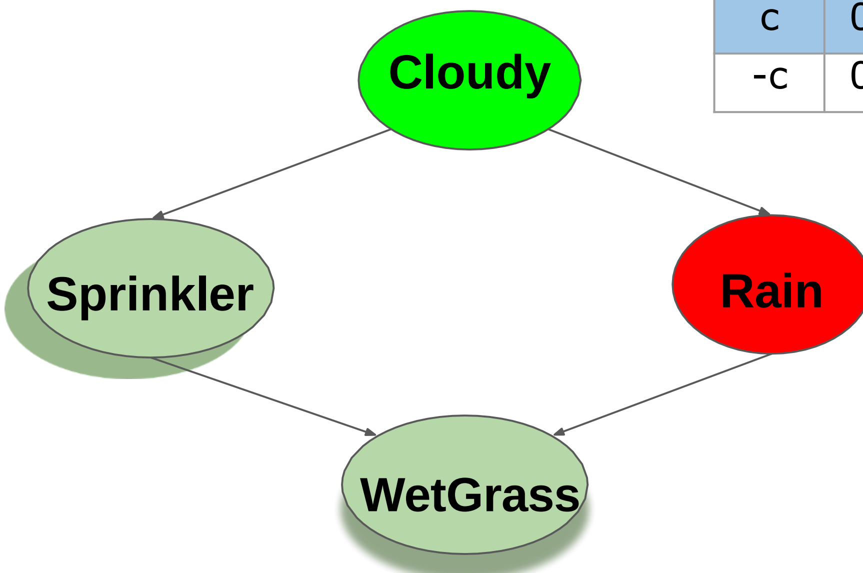

$P(S|C)$

c	s	0.1
	-s	0.9
-c	s	0.5
	-s	0.5

$P(R|C)$

c	r	0.8
	-r	0.2
-c	r	0.2
	-r	0.8

$P(W|S,R)$

s	r	w	0.99
		-w	0.01
	-r	w	0.90
		-w	0.10
-s	r	w	0.90
		-w	0.10
	-r	w	0.01
		-w	0.99

$u = 0.95$

- **Fix sampling order {C, S, R,W}**
- **Likelihood weighted sampling with (+s, +w)**
  - **{+c, +s, -r, _} , W = 1*0.1**

Example : Likelihood Weighting

Evidence: (+s,+w)

$P(C)$

c	0.5
-c	0.5

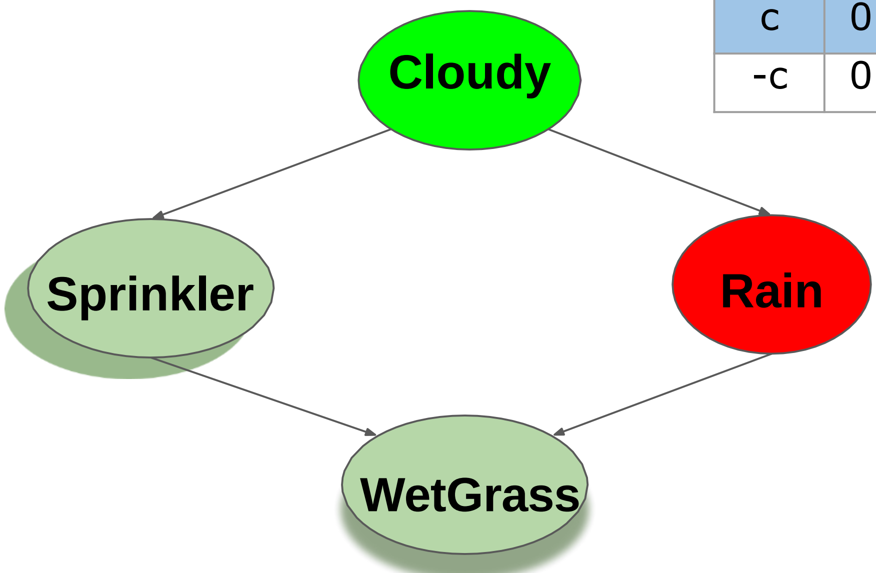

$P(S|C)$

c	s	0.1
	-s	0.9
-c	s	0.5
	-s	0.5

$P(R|C)$

c	r	0.8
	-r	0.2
-c	r	0.2
	-r	0.8

$P(W|S,R)$

s	r	w	0.99
		-w	0.01
	-r	w	0.90
		-w	0.10
-s	r	w	0.90
		-w	0.10
	-r	w	0.01
		-w	0.99

- **Fix sampling order {C, S, R,W}**
- **Likelihood weighted sampling with (+s, +w)**
  - **{+c, +s, -r, +w} , W = 1*0.1*0.9**

$u = 0.78$

Not Used

$P(+s|+c)$ $P(+w|+s,-r)$

Evidence: (+s,+w)

Example : Likelihood Weighting

Evidence: (+s,+w)

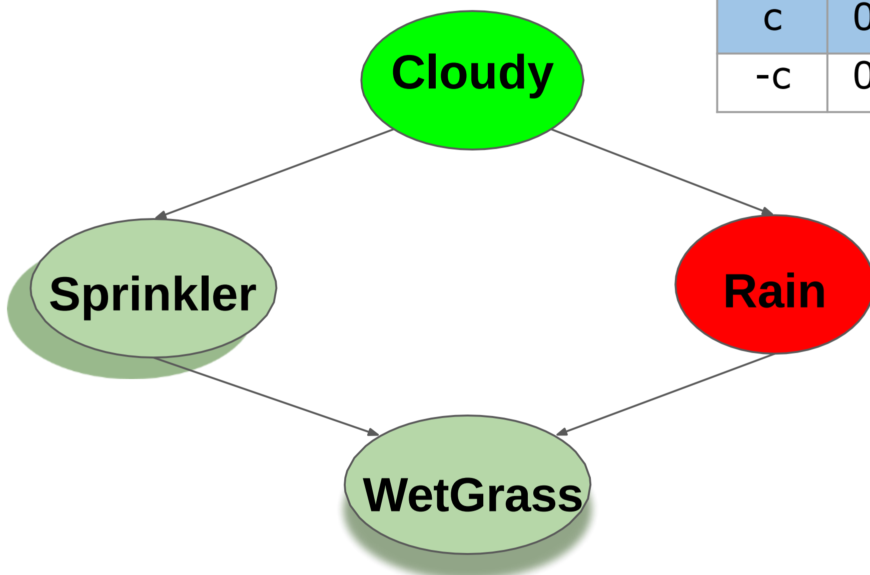

$P(C)$

c	0.5
-c	0.5

$P(S|C)$

c	s	0.1
	-s	0.9
-c	s	0.5
	-s	0.5

$P(R|C)$

c	r	0.8
	-r	0.2
-c	r	0.2
	-r	0.8

$P(W|S,R)$

s	r	w	0.99
		-w	0.01
	-r	w	0.90
		-w	0.10
-s	r	w	0.90
		-w	0.10
	-r	w	0.01
		-w	0.99

- **Fix sampling order {C, S, R,W}**
- **Likelihood weighted sampling with (+s, +w)**
  - **{+c, +s, -r, +w} , W = 0.09**

w is small,
meaning we have less confidence in our sample.

Example : Likelihood Weighting

Evidence: (+s,+w)

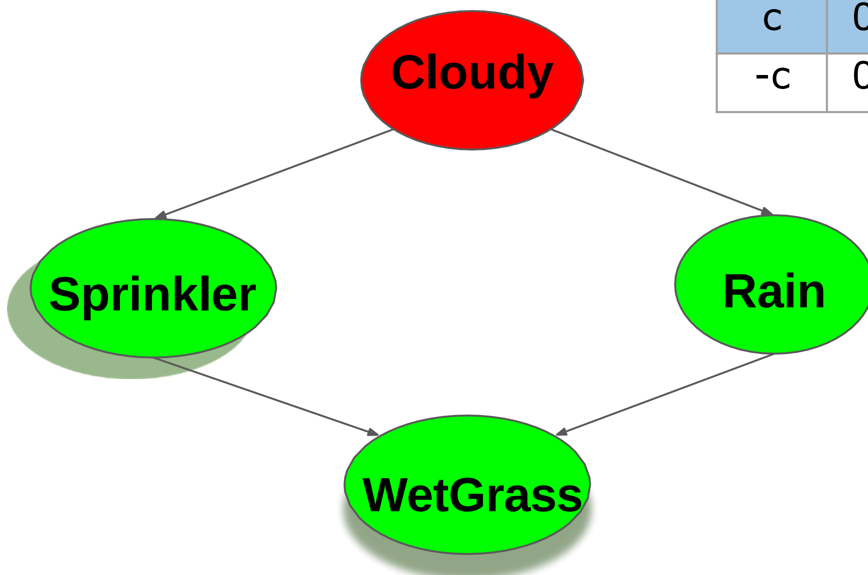

$P(C)$

c	0.5
-c	0.5

$P(S|C)$

c	s	0.1
	-s	0.9
-c	s	0.5
	-s	0.5

$P(R|C)$

c	r	0.8
	-r	0.2
-c	r	0.2
	-r	0.8

$P(W|S,R)$

s	r	w	0.99
		-w	0.01
	-r	w	0.90
		-w	0.10
-s	r	w	0.90
		-w	0.10
	-r	w	0.01
		-w	0.99

- **Fix sampling order {C, S, R,W}**
- **Likelihood weighted sampling with (+s, +w)**
  - **{+c, +s, -r, +w}, W = 0.09**
  - **{-c, +s, -r, +w}, W = 0.45**
  - **….**
  - **{+c, +s, +r, +w}, W = 0.099**

Example : Likelihood Weighting

Say the following $N = 100$ samples were generated from the wetgrass network with associated likelihood/weights:

c s r w

- 3 samples of $[true, true, true, true]$ with $w = 0.099$
- 2 samples of $[true, true, false, true]$ with $w = 0.09$
- 55 samples of $[false, true, true, true]$ with $w = 0.495$
- 40 samples of $[false, true, false, true]$ with $w = 0.45$

Notice that all samples are consistent with the evidence $Sprinkler = true$ and $WetGrass = true$.

The desired probability estimate is

P(R|s,w)=P(R,s,w)/P(s,w)= αP(R,s,w)= $\alpha \sum_{C} \boldsymbol{P(C, R, s, w,)}$

$$\hat{P}(Rain|Sprinkler = true, WetGrass = true)$$

(c,r,s,w) (-c,r,s,w)

$$= \alpha\langle 3 \times 0.099 + 55 \times 0.495, 2 \times 0.09 + 40 \times 0.45\rangle$$

(c,-r,s,w) (-c,-r,s,w)

$$= \alpha\langle 27.522, 18.18\rangle = \langle 0.60, 0.40\rangle$$

Likelihood Weighting

$$P(X|e)$$

For an arbitrary Bayesian Network, let X be the query variable, e be the values of the evidence variables and Y be the unobserved variables.

Let

- $N_{WS}(X, Y, e)$ be the number of samples generated for the event X, Y and e.
- $w(X, Y, e)$ be the weight of a sample corresponding to the event X, Y and e.

The the estimate for $P(X|e)$ is

$$\hat{P}(X|e) = \alpha \sum_{\forall Y} N_{WS}(X, Y, e) w(X, Y, e)$$

Likelihood Weighting : Why it works?

- In a BN, let **E** represents all evidence variables, **Z** represents all nonevidence variables including the query variable X. The sampling probability distribution is:

l: the number of nonevidence variables.

$$S_{WS}(\mathbf{z}, \mathbf{e}) = \prod_{i=1}^{l} P(z_i \,|\, parents(Z_i))$$

The sampling process: $S_{WS}(C, s, R, w) = P(C)P(R|C)$

- Now, samples have weights.

m: the number of nonevidence variables.

$$w(\mathbf{z}, \mathbf{e}) = \prod_{i=1}^{m} P(e_i \,|\, parents(E_i))$$

e.g., $w = P(s|c)P(w|s,-r)$ for {+c, +s, -r, +w}

Likelihood Weighting : Why it works?

- Together, weighted sampling distribution is consistent.

$$
\begin{aligned}
S_{WS}(\mathbf{z},\mathbf{e})w(\mathbf{z},\mathbf{e}) &= \prod_{i=1}^{l} P(z_i \mid parents(Z_i)) \prod_{i=1}^{m} P(e_i \mid parents(E_i)) \\
&= P(\mathbf{z},\mathbf{e})
\end{aligned}
$$

- $$
\begin{aligned}
P(X|e) &= \alpha \sum_{\forall Y} P(X,Y,e) \\
&= \alpha \sum_{\forall Y} S_{WS}(X,Y,e)w(X,Y,e) \\
&= \lim_{N\to\infty} \hat{\alpha} \sum_{\forall Y} N_{WS}(X,Y,e)w(X,Y,e) \\
&= \lim_{N\to\infty} \hat{P}(X|e)
\end{aligned}
$$

Recall $\hat{P}(X|e) = \hat{\alpha} \sum_{\forall Y} N_{WS}(X,Y,e)w(X,Y,e)$

Inference on Bayesian Networks

Approximate inference with sampling:

Also called Monte Carlo Methods.

- Direct Sampling (Prior Sampling).
- Rejection Sampling.
- Likelihood Weighting.
- Gibbs Sampling

Direct sampling

Markov chain sampling

Gibbs Sampling

- Gibbs sampling is a special instance of **Markov Chain Monte Carlo (MCMC) Methods.**

  Generates an event by making a random change to preceding event

  - Think that network is in a **current state** which species an event.
  - **Next state** is reached by sampling a value for one non-evidence variable X to be conditioned on the current values of X's **Markov blanket variables.**
  - Gibbs sampler thus wonders randomly in the state space by flipping one variable at a time while keeping evidence variables fixed.

Recall Markov Blanket

Recall that the Markov Blanket of a variable comprises of the parents, children, and children's parents of the variable.

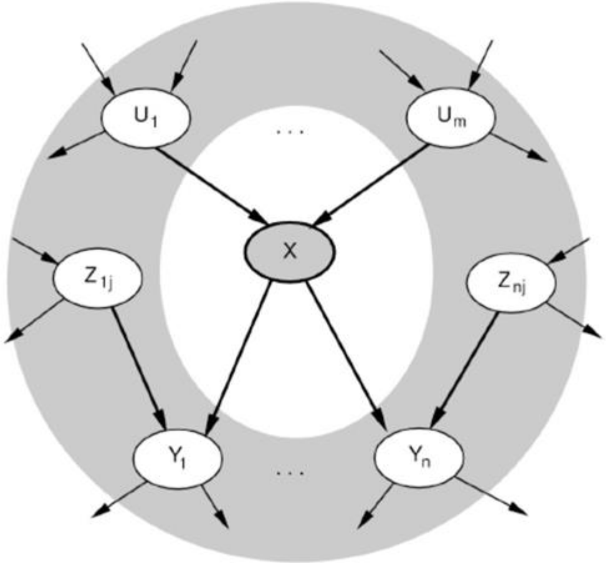

A node is conditionally independent of all others given the Markov Blanket of the node.

Gibbs Sampling Example (P(S|+r))

Step 1: initialize evidence

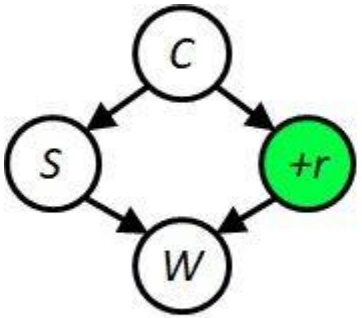

Gibbs Sampling Example (P(S|+r))

Step 1: initialize evidence

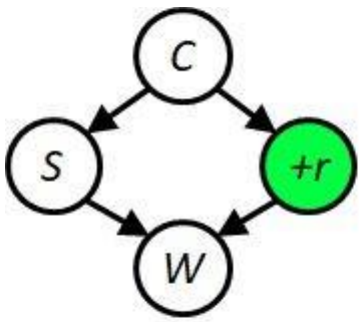

Step 2: initialize other variables (random)

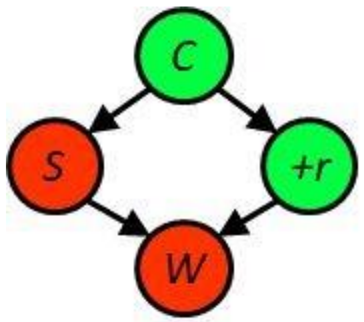

Initial state, e.g,: {+c, -s, +r, –w}

Gibbs Sampling Example (P(S|+r))

Step 1: initialize evidence

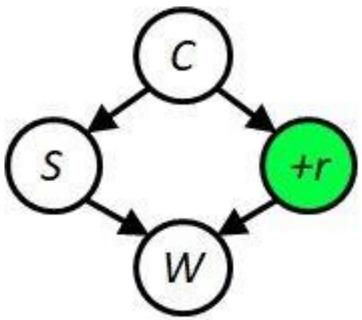

Step 2: initialize other variables (random)

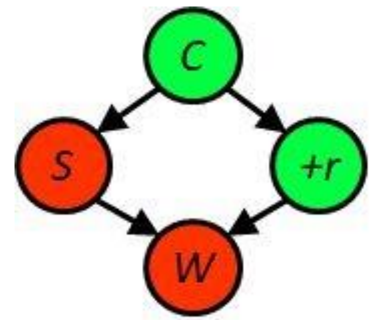

Initial state, e.g,: {+c, -s, +r, –w}

Step 3: Repeat following

- Choose a nonevidence variable X (at random). Here X is S, C, W.
- Sample X given the current values of X's Markov blanket variables.

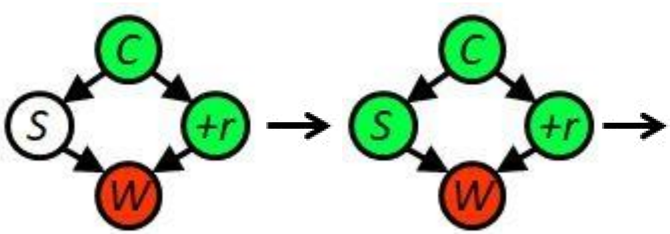

Sample from $P(S|+c, -w, +r)$

Suppose the result is true, the we get a new sample {+c, +s, +r, –w}

Gibbs Sampling Example (P(S|+r))

Step 1: initialize evidence

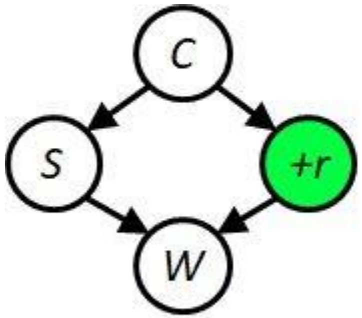

Step 2: initialize other variables (random)

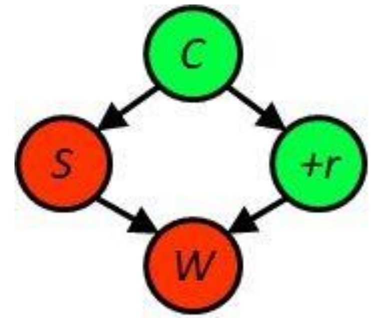

Initial state, e.g,: {+c, -s, +r, –w}

Step 3: Repeat following

- Choose a nonevidence variable X (at random). Here X can be S, C, W.
- Sample X given the current values of X's Markov blanket variables.

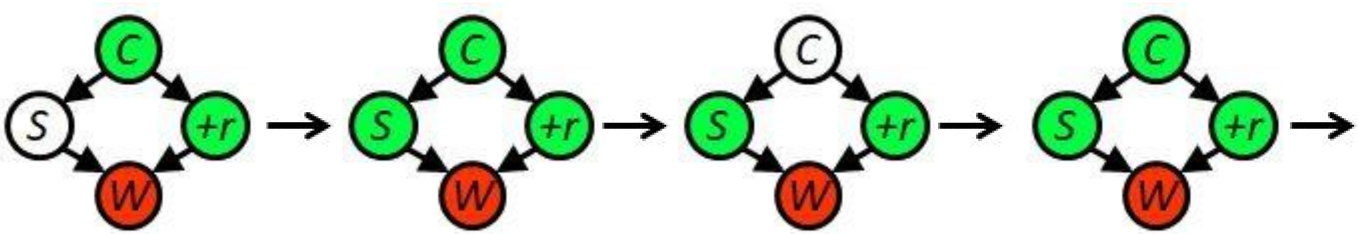

{+c, +s, +r, –w}

Sample from $P(C|+s,+r)$

Suppose the result is true, then we get another sample {+c, +s, +r, –w}

Gibbs Sampling Example (P(S|+r))

Step 1: initialize evidence

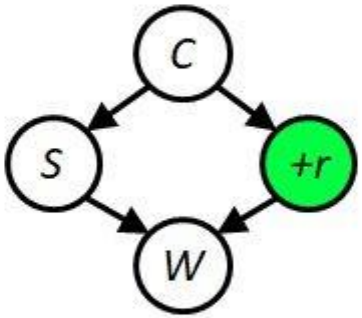

Step 2: initialize other variables (random)

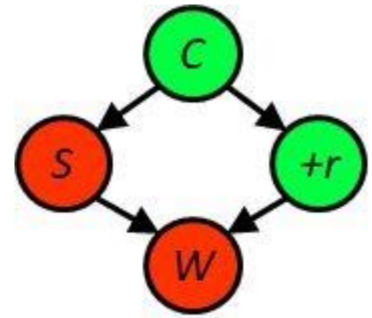

Initial state, e.g,: {+c, -s, +r, –w}

Step 3: Repeat following

- Choose a nonevidence variable X (at random). Here X can be S, C, W.
- Sample X given the current values of X's Markov blanket variables.

....

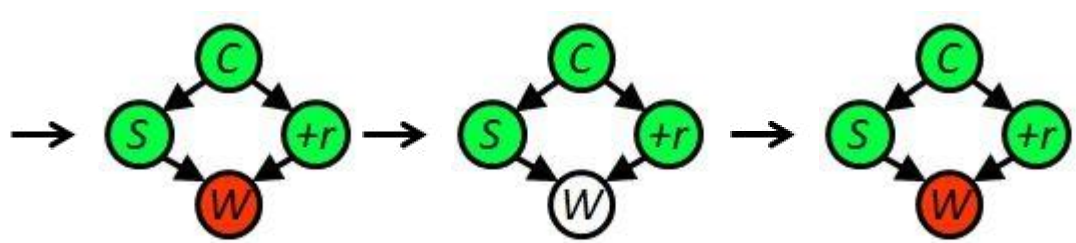

{+c, +s, +r, –w}

Sample from $P(W|+s, +r)$

{+c, +s, +r, -w}

Gibbs Sampling Example (P(S|+r))

Now suppose we get 100 samples with Gibbs Sampling.

- All samples were satisfying observation, i.e. {**Rain = true**}
- 37 of them had {**Sprinkler = true**}
- Which means, 63 of them had {**Sprinkler = false**}

$$P(S|Rain = true) = \alpha < 37, 63 >$$

$$= < 0.37, 0.63 >$$

Gibbs Sampling Example (P(S|+r))

$$\text{Sample from } P(S| + c, -w, +r)$$
$$\text{Sample from } P(C| + s, +r)$$

$$P\left(X|mb(X)\right)$$

- As in BN,

P(X|variable values of Markov Blanket of X)

$$= \alpha P\left(X|parents\left(X\right)\right) \times \prod_{Y_j \in Children(X)} P\left(Y_j|parents\left(Y_j\right)\right)$$

That is,

P(S|+c, -w, +r) = αP(S|+c)P(-w|+r)

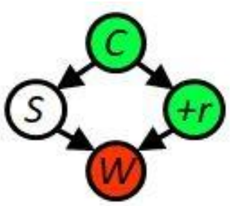

P(C|+s, +r) = α'P(C)P(+s|C)P(+r|C)